The Bi-LSTM's gap to LR closes as training data grow

−0.114 at 323K training windows, +0.042 at 1.12M. Grey: development folds (4-5 training years); blue: fresh test (all 8)

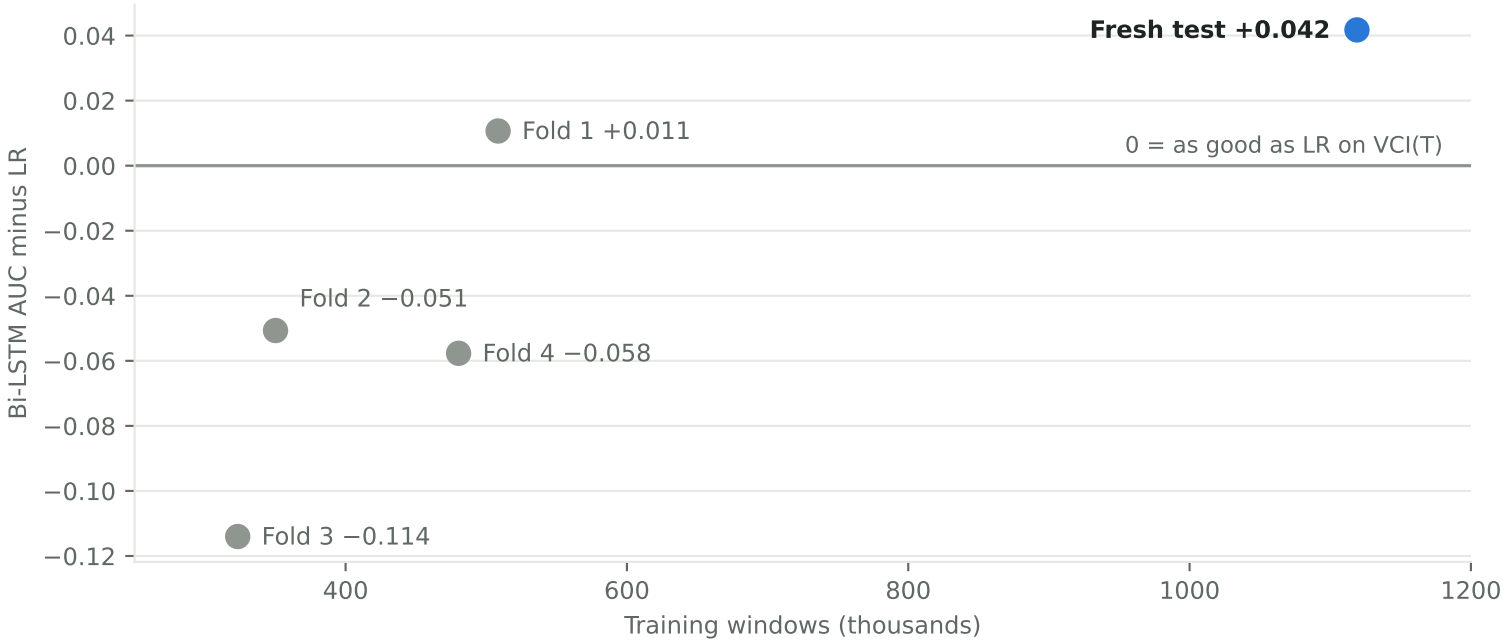

Source: results/scores/fresh_posthoc.json → dev_folds. Five points; training size moves with each fold's years and regions.